



HEXWARE

TRAINING PROJECT REPORT

INVENTORY MANAGEMENT

**OBJECTIVE : Analysis & Reporting
System of Inventory Management.**

PYTHON BATCH - 2

MEMBERS (TEAM - 14):

- **Rushmitha Sreeja Kaluvakolan**
- **Katika Mubeena**



Introduction

This mini-project focuses on developing an **Inventory Management Analysis and Reporting System** using **PySpark**. The project centers on building a data warehouse based on a star or snowflake schema, transforming OLTP source data into OLAP structures, applying data transformations, and generating insightful reports. By leveraging tools like Python and PySpark on the Databricks platform, the project integrates **Microsoft Azure Services** for efficient data storage, processing, and report generation.

The primary objective was to explore the synergy between **Big Data technologies** and **cloud services** for working with large datasets, gaining hands-on experience in implementing a robust data analysis pipeline.

Technologies Used

Programming and Frameworks

- **Python**
- **Apache Spark** (Version 15.4 LTS)

Microsoft Azure Services

1. **Azure Data Lake Storage**: Securely storing large volumes of data.
2. **Azure Databricks**: For data processing and analysis using:
 - Databricks Clusters
 - Databricks Catalog
 - Databricks Workspace
3. **Azure SQL Server**: Managing and querying structured data.
4. **SQL Database**: Facilitating operations on structured datasets.
5. **Azure Data Factory**: Automating data workflows and orchestrations.

This setup enabled efficient data analysis and processing while providing practical experience with cloud-based tools essential for modern data projects.



Abstract

In today's competitive business landscape, efficient inventory management is vital for optimizing operational performance and customer satisfaction. This project focuses on developing an **Inventory Management System** using dimensional modeling to enable seamless data analysis, reporting, and decision-making. The system is designed to handle and integrate large-scale transaction data with key dimensions, such as customers, products, sellers, and time, into a structured schema for business intelligence and analytics.

The project architecture comprises a **central fact table** that records transactional details, including sales, costs, and profits, linked to multiple **dimension tables**—Time, Customers, Products, Sellers, and Transactions—through carefully designed relationships. Data from various sources is ingested into an Azure-based ecosystem leveraging **Azure Databricks** for processing and transformation. The **Time Table** provides a comprehensive breakdown of dates, facilitating insights into sales trends across time dimensions, such as years, months, and quarters.

This system utilizes **PySpark** for data cleaning, transformation, and ETL (Extract, Transform, Load) processes, ensuring scalability and efficiency. Reports such as **top customers by transaction value**, **state-wise transaction summaries**, and **seller performance metrics** are generated and stored back into Azure Data Lake in accessible formats (CSV/Parquet). The project integrates robust error handling mechanisms to address missing data, failed relationships, and performance bottlenecks.

By implementing this inventory management system, organizations can achieve improved data visibility, generate actionable insights, and support strategic decision-making, ultimately enhancing overall business performance. The solution serves as a foundation for scalable analytics in inventory control and sales optimization.



DIMENSION TABLES

Customer Dimension

Fields in the Customer Dimension Table:

- CUSTOMER_ID
- Customer Name
- CUSTOMER_LOGIN_ID
- CUSTOMER_STREET_ADDRESS
- CUSTOMER_CITY
- CUSTOMER_STATE
- CUSTOMER_ZIP
- CUSTOMER_PHONE_NO.

Products Dimension

Fields in the Products Dimension Table:

- PRODUCT_ID
- CATEGORY_ID
- PRODUCT Name
- PRODUCT Brand
- Product Model No.
- PRODUCT_STOCK

SELLERS Dimension

Fields in the SELLERS Dimension Table:

- SELLER_ID
- SELLER_NAME



- SELLER_RATING
- SELLER_STREET_ADDRESS
- SELLER_CITY
- SELLER_STATE
- SELLER_ZIP
- SELLER_PHONE_NO.

TIME Dimension

Fields in the Time Dimension Table:

- date_key
- full_date
- day_of_week
- month_name
- quarter
- fb_year

TRANSACTIONS Dimension

Fields in the Transactions Dimension Table:

- TRANSACTION_ID
- TRANSACTION_DATE
- TRANSACTION_AMOUNT
- TRANSACTION_TYPE
- DISPATCH_DATE
- EXPECTED_DATE
- DELIVERY_DATE

FACT TABLE

Fields in the Fact Table:

- TRANSACTION_ID
- Date
- PRODUCT_ID
- CUSTOMER_ID
- SELLER_ID
- PRODUCT_COST_PRICE
- PRODUCT_SELLING_PRICE



- SELLER_ID
- SELLER_NAME
- SELLER_RATING
- SELLER_STREET_ADDRESS
- SELLER_CITY
- SELLER_STATE
- SELLER_ZIP
- SELLER_PHONE_NO.

TIME Dimension

Fields in the Time Dimension Table:

- date_key
- full_date
- day_of_week
- month_name
- quarter
- fb_year

TRANSACTIONS Dimension

Fields in the Transactions Dimension Table:

- TRANSACTION_ID
- TRANSACTION_DATE
- TRANSACTION_AMOUNT
- TRANSACTION_TYPE
- DISPATCH_DATE
- EXPECTED_DATE
- DELIVERY_DATE

REPORTS TO BE BUILT

- Transaction Amount Wise Top 10 Customers
- State wise number of orders
- Preferred mode of payment across states
- Quarterly profit for a particular year
- Quarterly sales count of each product category

Workflow: Azure Data Lake Storage Setup

To efficiently manage data, we implemented Azure Data Lake Storage and structured it into two organized containers:

Inventory

- This container holds the raw source data in CSV format, which forms the foundation of the project.
- It includes critical datasets such as transactions, customers, and products, used for data processing and transformation.
- These files serve as the primary input for performing analytics and building the ETL pipeline.

The screenshot displays the Microsoft Azure portal interface for the 'inventory33' storage account. The 'Containers' section is active, showing a list of containers. The table below represents the data visible in the screenshot:

Name	Last modified	Anonymous access level	Lease state
\$logs	12/2/2024, 12:03:23 PM	Private	Available
inventory	12/2/2024, 12:24:20 PM	Private	Available
reports	12/13/2024, 12:50:31 PM	Private	Available

Home > inventory33 | Containers >

inventory Container

Search

Upload Add Directory Refresh Rename Delete Change tier Acquire lease Break lease Give feedback

Overview

Diagnose and solve problems

Access Control (IAM)

Settings

Authentication method: Access key (Switch to Microsoft Entra user account)

Location: inventory

Search blobs by prefix (case-sensitive) Show deleted objects

Name	Modified	Access tier	Archive status	Blob type	Size	Lease state
[-] _azuretmpfolder\$	12/6/2024, 11:01:41
[-] Inventory_management.csv	12/13/2024, 12:57:55...					...
[-] inventory-output	12/6/2024, 11:08:26
[-] preferred_payment_mode...	12/13/2024, 12:58:39...					...
[-] sales_count.csv	12/13/2024, 12:58:43...					...
[-] statewide_ordercount.csv	12/13/2024, 12:58:37...					...
[-] top10_customers.csv	12/13/2024, 12:58:34...					...
[-] yearly_profit.csv	12/13/2024, 12:58:41...					...
[-] DIM.Date.Table.csv	12/5/2024, 12:00:24 ...	Hot (Inferred)		Block blob	778.16 KiB	Available
[-] Inventory_Management_S...	12/2/2024, 12:24:35 ...	Hot (Inferred)		Block blob	281.21 KiB	Available

Output:

This container stores the transformed and processed CSV files generated as output.

These output files are the results of our analysis and transformations.

Home > inventory33 | Containers >

inventory Container

Search

Upload Add Directory Refresh Rename Delete Change tier Acquire lease Break lease Give feedback

Overview

Diagnose and solve problems

Access Control (IAM)

Settings

Authentication method: Access key (Switch to Microsoft Entra user account)

Location: inventory / inventory-output

Search blobs by prefix (case-sensitive) Show deleted objects

Name	Modified	Access tier	Archive status	Blob type	Size	Lease state
[-] [-]						...
[-] customers_dim	12/13/2024, 12:58:11...					...
[-] fact_table	12/13/2024, 12:58:23...					...
[-] products_dim	12/13/2024, 12:58:13...					...
[-] seller_dim	12/13/2024, 12:58:16...					...
[-] transaction_dim	12/13/2024, 12:58:19...					...

Azure Databricks Workspace

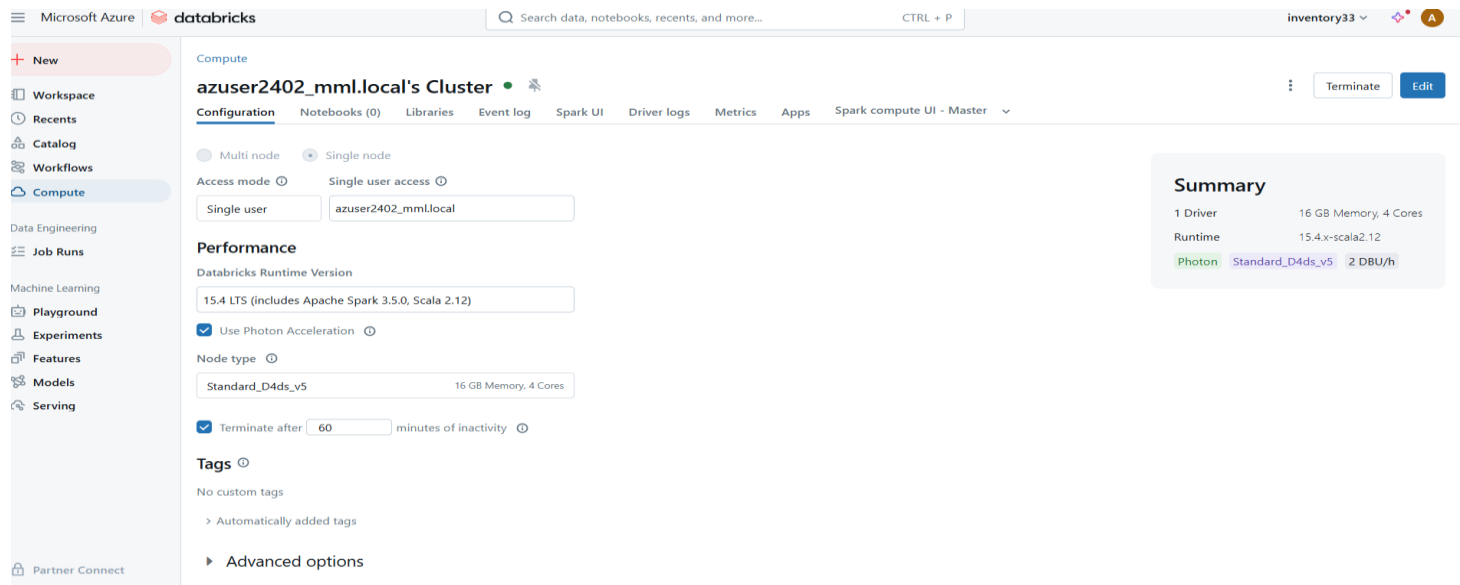
In the **Azure Databricks Workspace**, the following steps were performed:

1. Cluster Creation:

- 0 A **single-node cluster** was created to run and manage the notebook efficiently.
- 0 This cluster served as the processing environment for



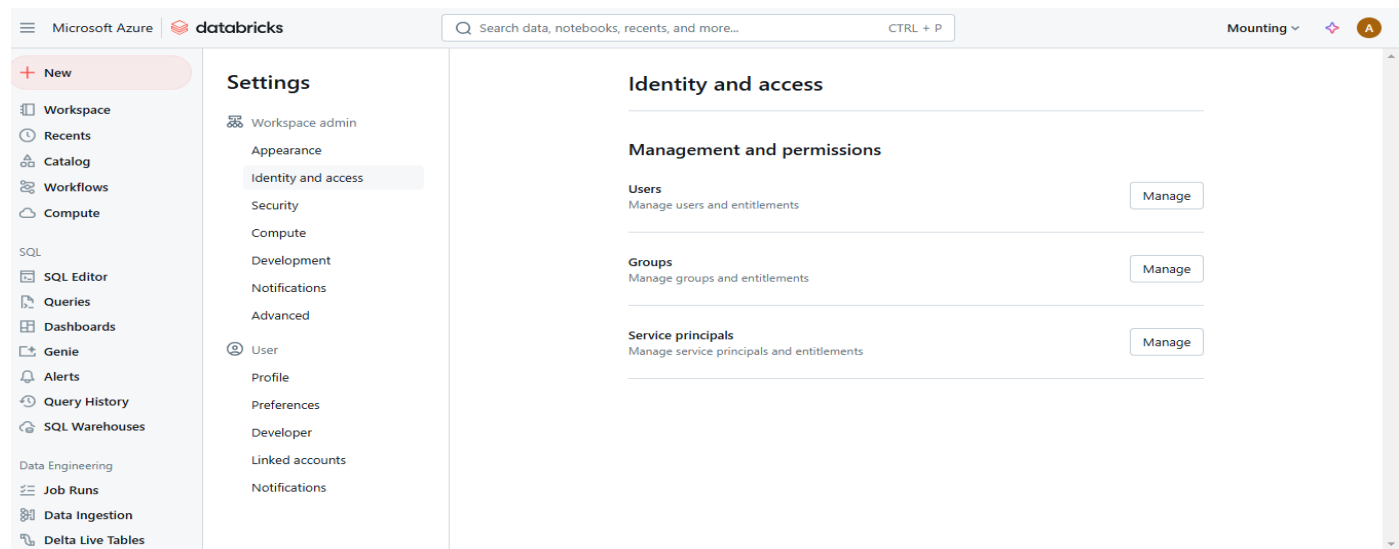
performing transformations and analysis on the data.



2. Notebook Organization:

A dedicated folder was created within the workspace to store all the notebooks.

These notebooks contained the PySpark code and other scripts used for dataprocessing and visualization.



The screenshot shows the Databricks web interface. On the left is a sidebar with navigation options: New, Workspace, Recents, Catalog, Workflows, Compute, Data Engineering, Job Runs, Machine Learning, Playground, Experiments, Features, Models, and Serving. The main area displays the 'Workspace' view for the user 'azuser2402_mml.local@techademy.com'. A table lists the contents of the workspace:

Name	Type	Owner	Created at
Inventory tables Notebook 2024-12-02 12...	Notebook	azuser2402_mml.local	2024-12-02 12:22:04
Inventory_Management (1)	Notebook	azuser2402_mml.local	2024-12-06 10:53:05

Mounting the Azure Data Lake Storage (ADLS) Account

To facilitate seamless access to the data stored in Azure Data Lake Storage (ADLS), the storage containers were mounted to the Databricks File System (DBFS) using the Account Key authentication method.

Mounting Containers

Input Container

1. Serves as the source for accessing raw data files, such as CSV files, required for transformations and data processing.
2. Enables Databricks to efficiently read and process unstructured and semi-structured data.

Output Container

1. Acts as the destination for storing processed and transformed data.
2. Includes results like reports, aggregated datasets, and intermediate outputs, ensuring data is well-organized for further use.

By mounting these containers, the integration between Databricks and ADLS ensures a smooth and efficient data pipeline for transformation and analysis

Mount input container(Inventor)

```

12:57 PM (11s) 4 Python
dbutils.fs.mount(
  source = "wasbs://inventory@inventory33.blob.core.windows.net",
  mount_point = "/mnt/inventory",
  extra_configs = {"fs.azure.account.key.inventory33.blob.core.windows.net": "OyFNfd0A5TkzgjbBgldfqqEXEpAdnDYK6zTGxIZ1sebVfjY4lCTeaZhkh4d2sa9kNqMgZyU0qDR9+ASth4M68A=="})
True

```

Mounting Container From ADLS for Output

```

12:57 PM (11s) 6
dbutils.fs.mount(
  source="wasbs://inventory@inventory33.blob.core.windows.net/",
  mount_point="/mnt/reports",
  extra_configs={"fs.azure.account.key.inventory33.blob.core.windows.net": "OyFNfd0A5TkzgjbBgldfqqEXEpAdnDYK6zTGxIZ1sebVfjY4lCTeaZhkh4d2sa9kNqMgZyU0qDR9+ASth4M68A=="})

```

PRE-PROCESSING

Read the data from Mount Point

```

12:57 PM (9s) 9 Python
df = spark.read.csv('dbfs:/mnt/inventory/Inventory_Management_Src1.csv', header=True, inferSchema=True)
(2) Spark Jobs
df: pyspark.sql.dataframe.DataFrame = [Date: string, TRANSACTION_ID: integer ... 30 more fields]

12:57 PM (4s) 10
df.write.format('parquet').mode('overwrite').save('/mnt/inventory/Inventory_management.csv',header=True)
(1) Spark Jobs

12:57 PM (1s) 11
df.show()
(1) Spark Jobs

```

Correcting the date format

```
12:57 PM (<1s) 26
from pyspark.sql.functions import col, year, quarter, month, dayofweek, monotonically_increasing_id

12:57 PM (1s) 27 Python
df2 = spark.read.csv('dbfs:/mnt/inventory/DIM.Date.Table.csv', header=True, inferSchema=True)

(2) Spark Jobs
df2: pyspark.sql.dataframe.DataFrame = [date_key: integer, full_date: timestamp ... 22 more fields]
```

DISPLAYING THE INPUT DATA

```

▶ 12:57 PM (2s) 17 Python
customers_dim.show(5)
▶ (2) Spark Jobs
+-----+-----+-----+-----+-----+-----+-----+
|CustomerID|CustomerName|CustomerLoginID|CustomerStreetAddress|CustomerCity|CustomerState|CustomerZip|CustomerPhoneNo|
+-----+-----+-----+-----+-----+-----+-----+
|100088|Gisela Mayer|5088|P.O. Box 321, 480...|SAN DIEGO|CALIFORNIA|60330|023-479-2802|
|100068|Tana Blevins|5068|Ap #364-8638 Et S...|BOSTON|MASSACHUSETTS|91583|081-828-4143|
|100074|Carly Parker|5074|9331 Ipsum. Ave|NEW YORK|NEWYORK|66948|016-408-0204|
|100003|Fritz Grant|5003|Ap #897-7736 Eges...|PHILADELPHIA|PENNSYLVANIA|65342|029-197-3614|
|100042|Leroy Pugh|5042|6658 Venenatis Rd.|AUSTIN|TEXAS|63417|050-412-9084|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

```

```

▶ 12:57 PM (<1s) 21
products_dim.show(5)
▶ (2) Spark Jobs
+-----+-----+-----+-----+-----+-----+
|ProductID|CategoryID|ProductName|ProductBrand|ProductModelNo|ProductStock|
+-----+-----+-----+-----+-----+-----+
|3000|4000|CAMERA|NIKON|1100|10|
|3089|4006|PENDRIVE|TRANSCEND|9976|9|
|3098|4002|BAG|AMERICAN TOURISTER|37950|9|
|3042|4007|EARPHONES|SOUND MAGIC|11540|9|
|3028|4009|SUNGLASSES|GUCCI|6270|9|
+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

```

```

▶ 12:57 PM (<1s) 25
seller_dim.show(5)
▶ (2) Spark Jobs
+-----+-----+-----+-----+-----+-----+-----+
|SellerID|SellerName|SellerRating|SellerStreetAddress|SellerCity|SellerState|SellerZip|SellerPhoneNo|
+-----+-----+-----+-----+-----+-----+-----+
|200049|Cherokee Richard|5|Ap #381-3851 Eget...|PARIS|ILE-DE-FRANCE|72816|01 33 96 49 60|
|200031|Robert Becker|1|182-4566 Gravida ...|ANGERS|PAYS DE LA LOIRE|90358|01 39 47 76 38|
|200001|Aaron Cooper|2|P.O. Box 446, 302...|LE MANS|PAYS DE LA LOIRE|94302|07 25 54 31 55|
|200021|Bo Page|3|Ap #257-6739 Lect...|GRENOBLE|RHONE-ALPES|65860|03 18 25 76 17|
|200096|Mary Lloyd|1|6075 Mus. Av.|REIMS|CHAMPAGNE-ARDENNE|80672|01 41 65 21 82|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

```



time_date_dim.show()

(1) Spark Jobs

YEAR	QUARTER	MONTH	DAY_OF_WEEK	DATE	TimeID
2008	1	January	2	20080101	0
2008	1	January	3	20080102	1
2008	1	January	4	20080103	2
2008	1	January	5	20080104	3
2008	1	January	6	20080105	4
2008	1	January	7	20080106	5
2008	1	January	1	20080107	6
2008	1	January	2	20080108	7
2008	1	January	3	20080109	8
2008	1	January	4	20080110	9
2008	1	January	5	20080111	10
2008	1	January	6	20080112	11
2008	1	January	7	20080113	12
2008	1	January	1	20080114	13
2008	1	January	2	20080115	14
2008	1	January	3	20080116	15

CREATING FACT AND DIMENSIONAL TABLES

Creating Seller DIM table

```

seller_dim = df11.select(
    "SELLER_ID",
    "SELLER_NAME",
    "SELLER_RATING",
    "SELLER_STREET_ADDRESS",
    "SELLER_CITY",
    "SELLER_STATE",
    "SELLER_ZIP",
    "SELLER_PHONE_NO"
)

```

seller_dim: pyspark.sql.dataframe.DataFrame = [SELLER_ID: integer, SELLER_NAME: string ... 6 more fields]

```

seller_dim = seller_dim.withColumnRenamed("SELLER_ID", "SellerID") \
    .withColumnRenamed("SELLER_NAME", "SellerName") \
    .withColumnRenamed("SELLER_RATING", "SellerRating") \
    .withColumnRenamed("SELLER_STREET_ADDRESS", "SellerStreetAddress") \
    .withColumnRenamed("SELLER_CITY", "SellerCity") \
    .withColumnRenamed("SELLER_STATE", "SellerState") \
    .withColumnRenamed("SELLER_ZIP", "SellerZip") \
    .withColumnRenamed("SELLER_PHONE_NO", "SellerPhoneNo")

```

seller_dim: pyspark.sql.dataframe.DataFrame = [SellerID: integer, SellerName: string ... 6 more fields]

creating dim table

```

transaction_dim = df11.select(
    "TRANSACTION_ID",
    "TRANSACTION_DATE",
    "TRANSACTION_AMOUNT",
    "TRANSACTION_TYPE",
    "DISPATCH_DATE",
    "EXPECTED_DATE",
    "DELIVERY_DATE"
)

```

transaction_dim: pyspark.sql.dataframe.DataFrame = [TRANSACTION_ID: integer, TRANSACTION_DATE: string ... 5 more fields]

```

transaction_dim = transaction_dim.withColumnRenamed("TRANSACTION_ID", "TransactionID") \
    .withColumnRenamed("TRANSACTION_DATE", "TransactionDate") \
    .withColumnRenamed("TRANSACTION_AMOUNT", "TransactionAmount") \
    .withColumnRenamed("TRANSACTION_TYPE", "TransactionType") \
    .withColumnRenamed("DISPATCH_DATE", "DispatchDate") \
    .withColumnRenamed("EXPECTED_DATE", "ExpectedDate") \
    .withColumnRenamed("DELIVERY_DATE", "DeliveryDate")

```

transaction_dim: pyspark.sql.dataframe.DataFrame = [TransactionID: integer, TransactionDate: string ... 5 more fields]

Inventory Management

```
output
```

```
seller_dim.show(5)
```

SellerID	SellerName	SellerRating	SellerStreetAddress	SellerCity	SellerState	SellerZip	SellerPhoneNo
200049	Cherokee Richard	5	Ap #381-3851 Eget...	PARIS	ILE-DE-FRANCE	72816 01 33 96 49 60	
200031	Robert Becker	1	182-4566 Gravida ...	ANGERS	PAYS DE LA LOIRE	90358 01 39 47 76 38	
200001	Aaron Cooper	2	P.O. Box 446, 302...	LE MANS	PAYS DE LA LOIRE	94302 07 25 54 31 55	
200021	Bo Page	3	Ap #257-6739 Lect...	GRENOBLE	RHONE-ALPES	65860 03 18 25 76 17	
200096	Mary Lloyd	1	6075 Mus. Av.	REIMS	CHAMPAGNE-ARDENNE	80672 01 41 65 21 82	

only showing top 5 rows

creating time DIM table

```
time_date_dim = df2.select(
    "fb_year",
    "QUARTER",
    "month_name",
    "DAY_OF_WEEK",
    "DATE_KEY"
)
```

```
time_date_dim = time_date_dim.withColumn("TimeID", monotonically_increasing_id())
```

```
time_date_dim = time_date_dim.withColumnRenamed("fb_year", "YEAR") \
    .withColumnRenamed("Quarter", "QUARTER") \
    .withColumnRenamed("Month_name", "MONTH") \
    .withColumnRenamed("Day_of_Week", "DAY_OF_WEEK") \
    .withColumnRenamed("Date_key", "DATE")
```

Output

```
time_date_dim.show()
```

[2008]	1	January	4	20080103	2
[2008]	1	January	5	20080104	3
[2008]	1	January	6	20080105	4
[2008]	1	January	7	20080106	5
[2008]	1	January	1	20080107	6
[2008]	1	January	2	20080108	7
[2008]	1	January	3	20080109	8
[2008]	1	January	4	20080110	9
[2008]	1	January	5	20080111	10
[2008]	1	January	6	20080112	11
[2008]	1	January	7	20080113	12
[2008]	1	January	1	20080114	13
[2008]	1	January	2	20080115	14
[2008]	1	January	3	20080116	15
[2008]	1	January	4	20080117	16
[2008]	1	January	5	20080118	17
[2008]	1	January	6	20080119	18
[2008]	1	January	7	20080120	19

only showing top 20 rows



Inventory Management

```
products_dim = products_dim.dropDuplicates(["ProductID"])
products_dim: pyspark.sql.dataframe.DataFrame = [ProductID: integer, CategoryID: integer ... 4 more fields]
```

Output

```
products_dim.show(5)
```

(2) Spark Jobs

ProductID	CategoryID	ProductName	ProductBrand	ProductModelNo	ProductStock
3000	4000	CAMERA	NIKON	1100	10
3089	4006	PENDRIVE	TRANSCEND	9976	9
3098	4002	BAG	AMERICAN TOURISTER	37950	9
3042	4007	EARPHONES	SOUND MAGIC	11540	9
3028	4009	SUNGLASSES	GUCCI	6270	9

only showing top 5 rows

Creating the Fact Table

```
inventory_fact = df11.select(
    "TRANSACTION_ID",
    "PRODUCT_ID",
    "CUSTOMER_ID",
    "SELLER_ID",
    "TRANSACTION_DATE",
    "PRODUCT_COST_PRICE",
    "PRODUCT_SELLING_PRICE"
)
```

inventory_fact: pyspark.sql.dataframe.DataFrame = [TRANSACTION_ID: integer, PRODUCT_ID: integer ... 5 more fields]

```
inventory_fact = inventory_fact.join(
    time_dim.select("TimeID", "Date"),
    inventory_fact.TRANSACTION_DATE == time_dim.Date,
    "left"
).drop("Date", "TRANSACTION_DATE")
```

inventory_fact: pyspark.sql.dataframe.DataFrame = [TRANSACTION_ID: integer, PRODUCT_ID: integer ... 5 more fields]

```
inventory_fact = inventory_fact.withColumn("FactID", monotonically_increasing_id())
```

inventory_fact: pyspark.sql.dataframe.DataFrame = [TRANSACTION_ID: integer, PRODUCT_ID: integer ... 6 more fields]

```
inventory_fact = inventory_fact.select(
    "FactID", "TimeID", "TRANSACTION_ID", "PRODUCT_ID",
    "CUSTOMER_ID", "SELLER_ID",
    "PRODUCT_COST_PRICE", "PRODUCT_SELLING_PRICE"
)
```

inventory_fact: pyspark.sql.dataframe.DataFrame = [FactID: long, TimeID: long ... 6 more fields]

```
inventory_fact.show(5)
```

(4) Spark Jobs

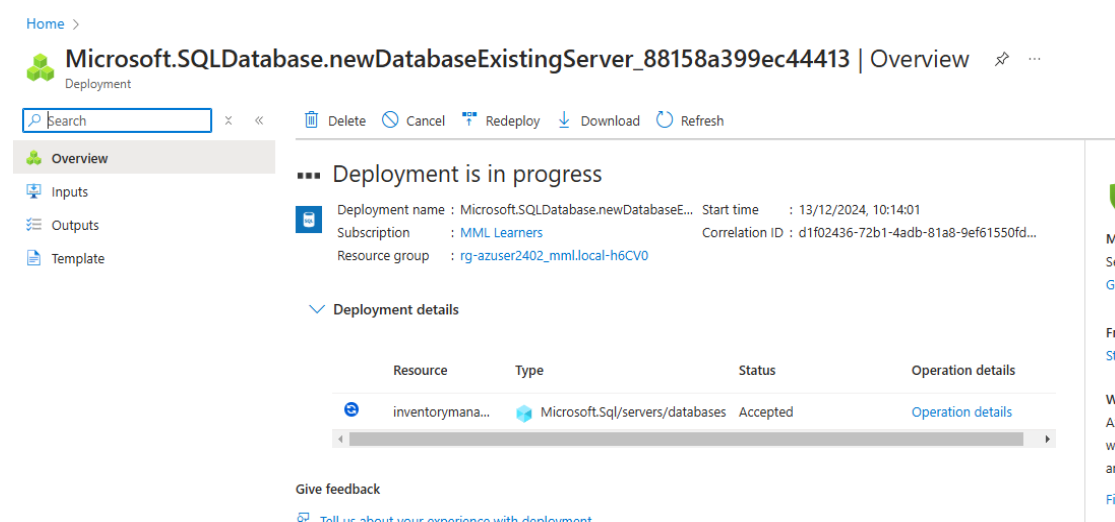
FactID	TimeID	TRANSACTION_ID	PRODUCT_ID	CUSTOMER_ID	SELLER_ID	PRODUCT_COST_PRICE	PRODUCT_SELLING_PRICE
0	734	5000	3100	100060	200074	32000	35000
1	734	5001	3001	100038	200091	2450	2450
2	734	5002	3012	100045	200048	3500	3500
3	322	5003	3089	100090	200055	600	650
4	875	5004	3023	100097	200025	22000	21500



STORING IN SQL DATABASE

First We created an Azure SQL Database Server

Azure SQL Server is a logical server that acts as a container for hosting multiple Azure SQL Databases. It provides the foundation for managing and scaling your relational databases in the cloud.



The screenshot shows the Azure Portal interface for a deployment. The breadcrumb navigation at the top reads "Home > Microsoft.SQLDatabase.newDatabaseExistingServer_88158a399ec44413 | Overview". The deployment name is "Microsoft.SQLDatabase.newDatabaseExistingServer_88158a399ec44413". The deployment status is "Deployment is in progress". The deployment details show the following information:

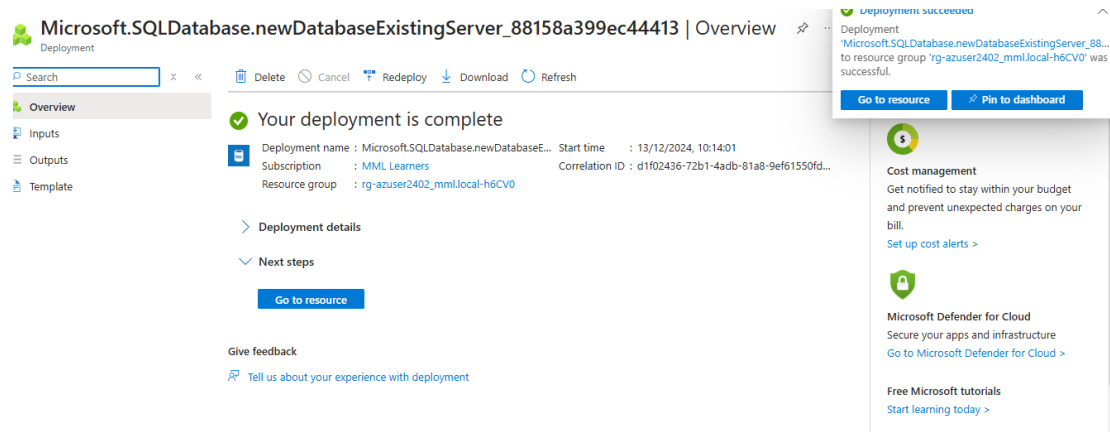
Resource	Type	Status	Operation details
inventorymana...	Microsoft.Sql/servers/databases	Accepted	Operation details

The deployment details also show the following information:

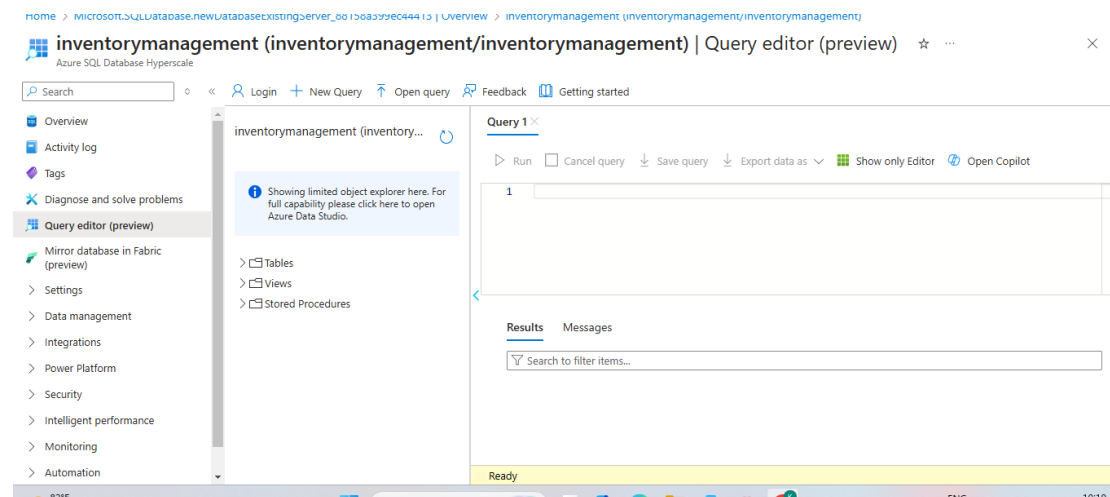
- Deployment name : Microsoft.SQLDatabase.newDatabaseExistingServer_88158a399ec44413
- Subscription : MML Learners
- Resource group : rg-azuser2402_mml-local-h6CV0
- Start time : 13/12/2024, 10:14:01
- Correlation ID : d1f02436-72b1-4adb-81a8-9ef61550fd...

Creating SQL Database

Azure SQL Database is a fully managed, relational database service built on Microsoft's SQL Server technology, designed for scalability, performance, and availability in the cloud. It supports a wide range of workloads, from small-scale applications to large-scale enterprise systems, making it an ideal choice for modern application development.



Storing all Dimensions and Fact Tables in Database



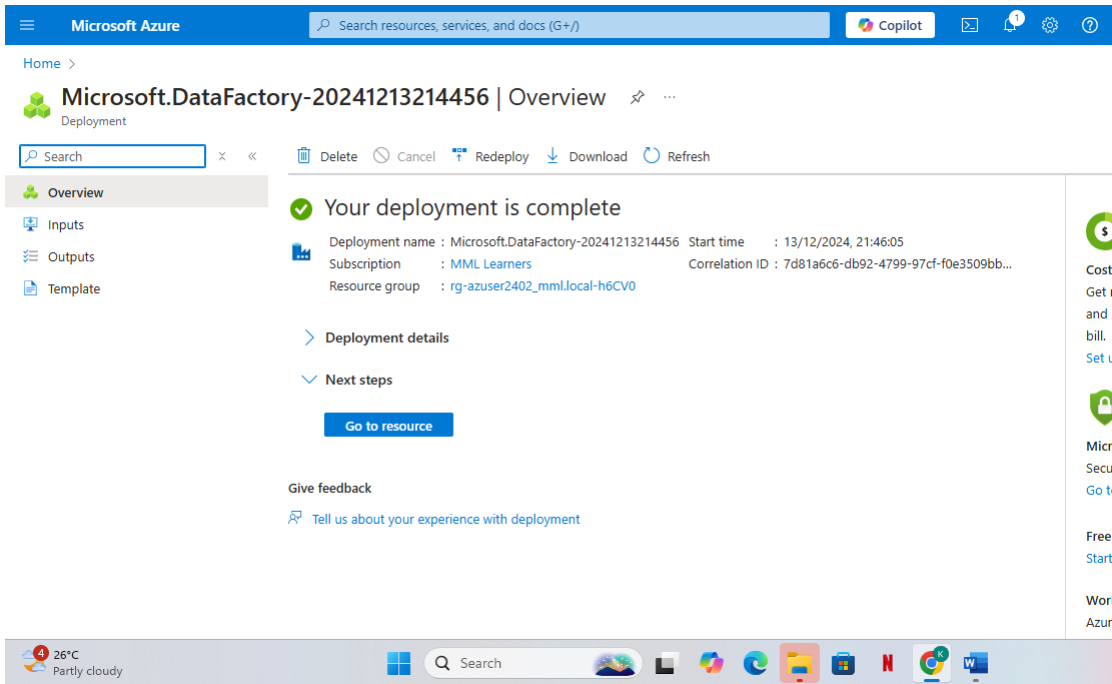
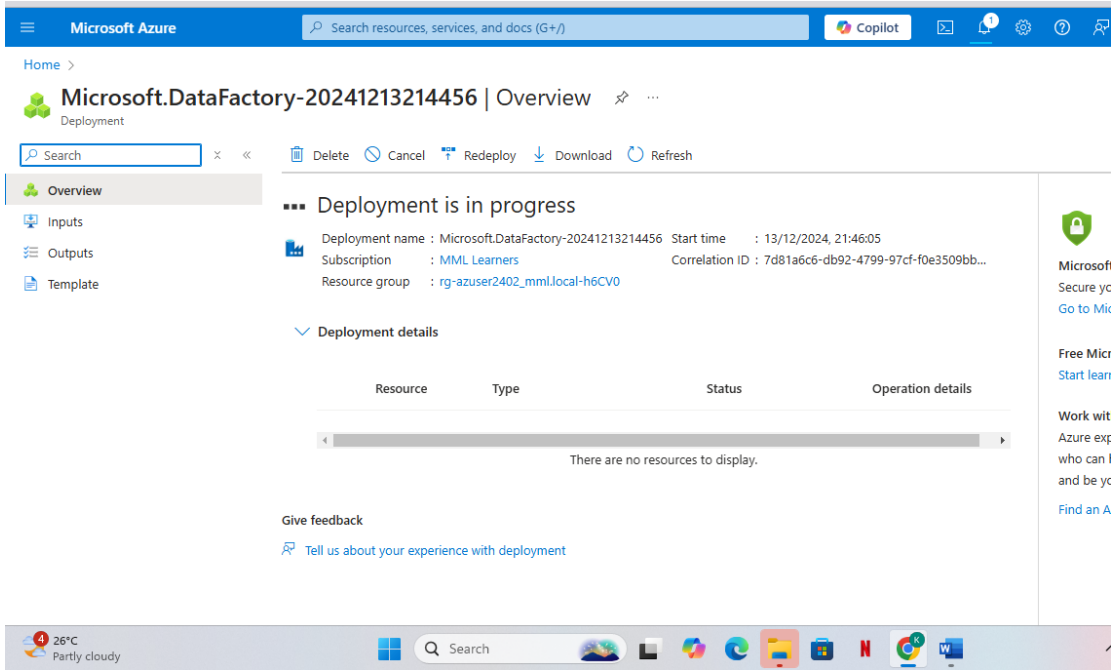
Azure Data Factory (ADF)

Azure Data Factory (ADF) is a cloud-based data integration and orchestration service provided by Microsoft Azure. It enables the creation, scheduling, and monitoring of data pipelines to efficiently transform and move data between diverse sources and destinations.

Key Features of Azure Data Factory:

- **Seamless Data Movement:** Facilitates data integration across on-premises, cloud, and hybrid environments.
- **ETL and ELT Pipelines:** Supports scalable and efficient **Extract, Transform, Load (ETL)** and **Extract, Load, Transform (ELT)** workflows, empowering organizations to process large volumes of data with ease.
- **Automation and Monitoring:** Provides tools for scheduling and monitoring data pipelines to ensure reliable and automated workflows.

- **Flexibility:** Supports a wide range of data formats, connectors, and transformations, making it a versatile solution for modern data workflows.

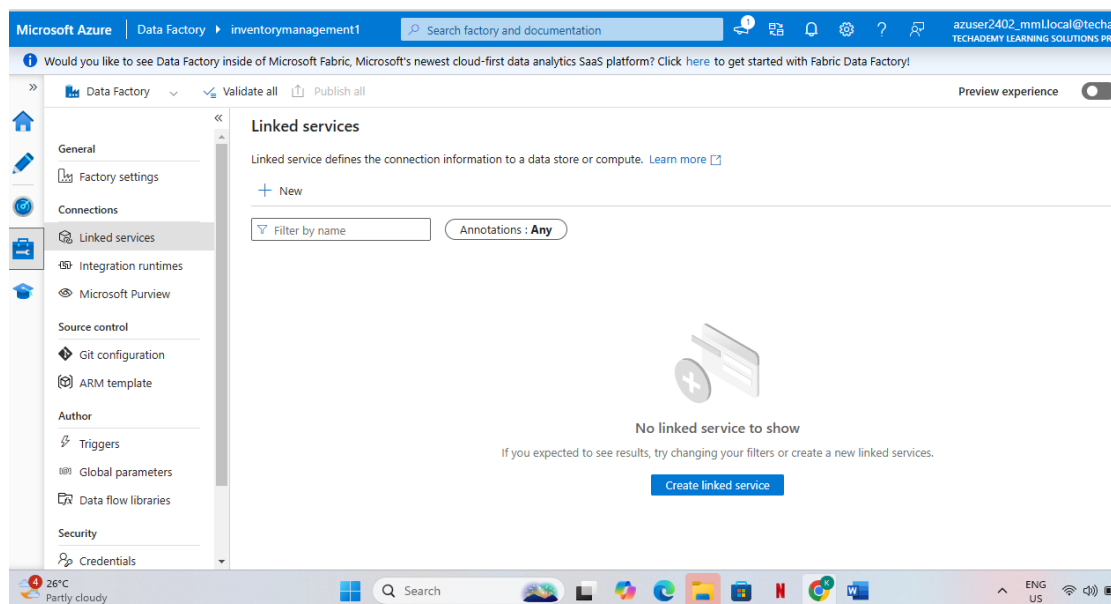


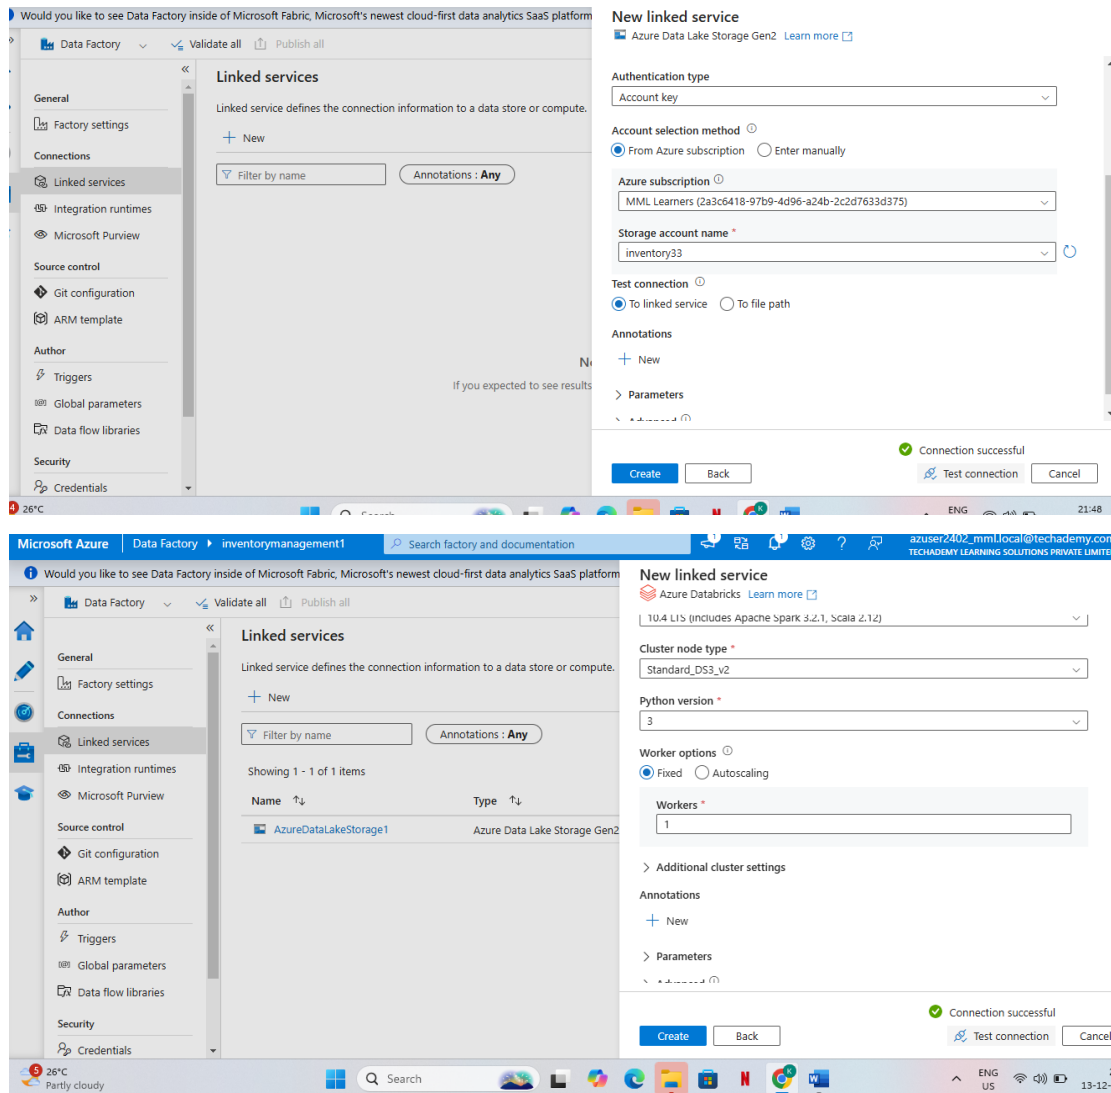
Creating an Azure Data Lake Storage (ADLS) Service Connection with Azure Data Factory (ADF)

Azure Data Lake Storage (ADLS) is a high-performance, enterprise-grade storage solution optimized for big data analytics. To enable Azure Data Factory (ADF) to interact with data stored in ADLS, a **service connection (linked service)** needs to be created.

This connection serves as a secure bridge, allowing ADF to seamlessly read, write, and manage data in ADLS as part of data integration and ETL workflows. By establishing this connection, ADF pipelines can efficiently perform tasks such as:

- **Data Ingestion:** Importing raw data from various sources into ADLS.
- **Data Transformation:** Processing and transforming data for analytical or operational purposes.
- **Data Export:** Moving processed data from ADLS to downstream systems or destinations





Creating Database Service Connection in Azure Data Factory (ADF)

Establishing a database service connection in Azure Data Factory (ADF) enables seamless integration with relational and non-relational databases. This connection allows ADF to move, transform, and process data effectively. In ADF, **service connections (linked services)** serve as connectors to external data sources, empowering pipelines to interact with databases hosted on Azure (e.g., Azure SQL Database, SQL Server), on-premises, or third-party cloud environments.

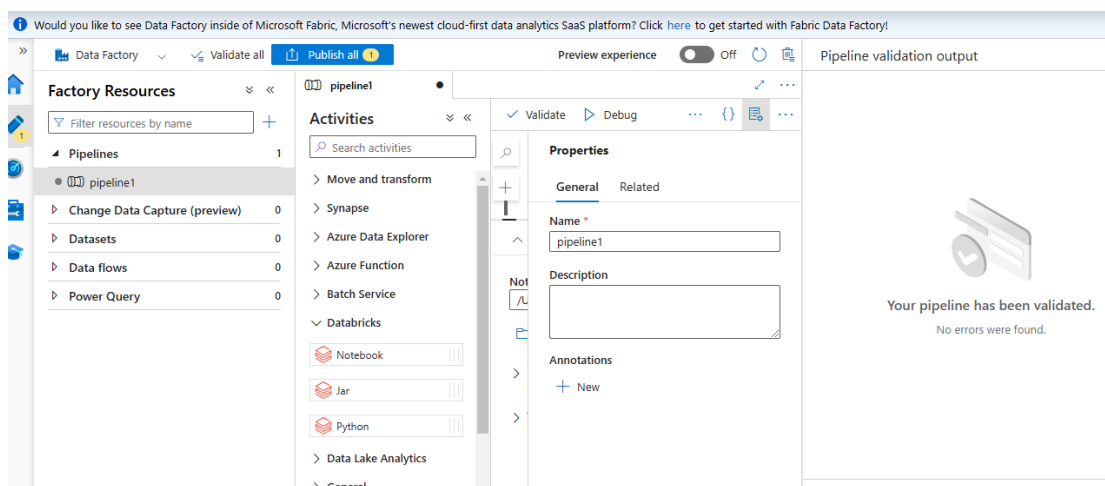
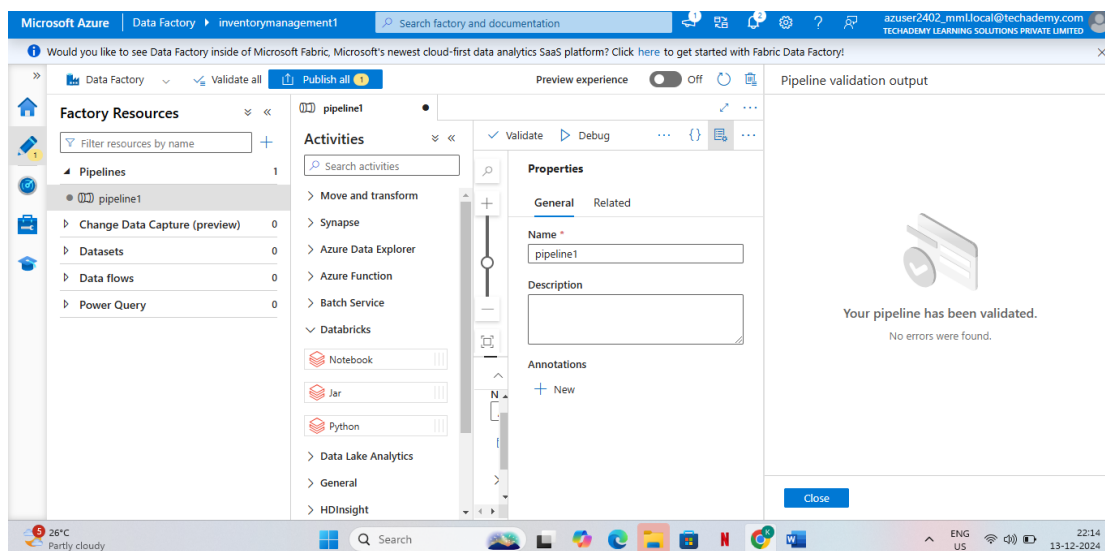
Service Connection:

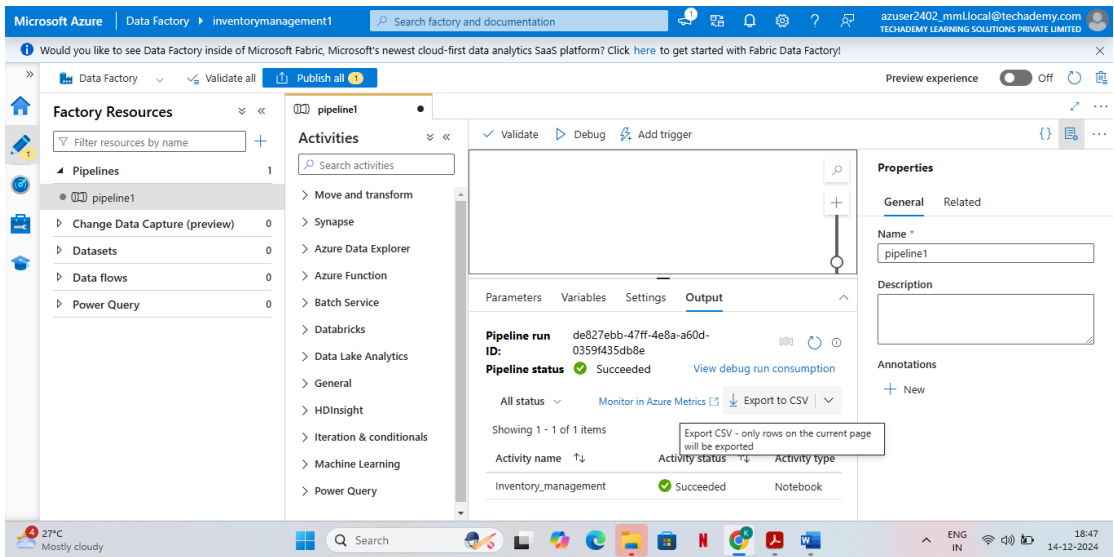
1. Provides ADF with the credentials and configuration to securely connect to database instances.
2. Enables ADF to perform read and write operations as part of ETL (Extract, Transform, Load) or ELT (Extract, Load, Transform) workflows.

Integration Runtime (IR):

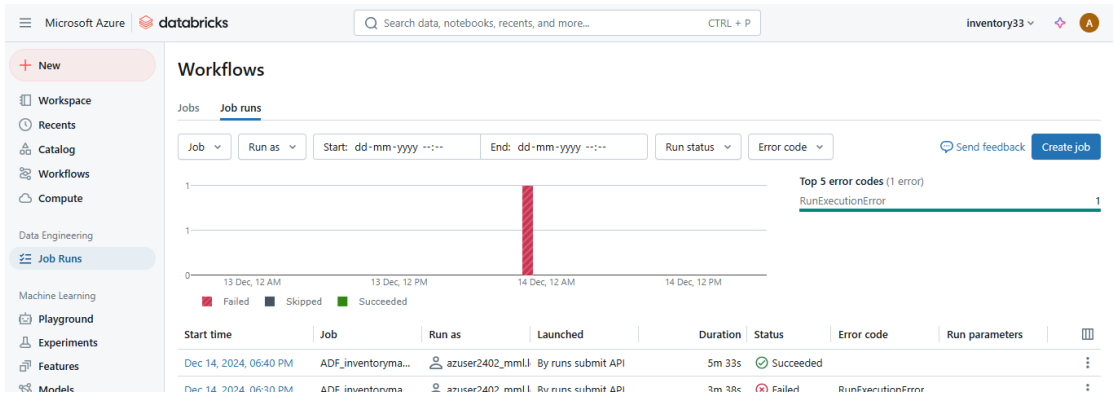
1. Acts as the compute infrastructure to move and transform data across cloud and on-premises data stores.
2. Facilitates secure and scalable data transfer and transformation for database connections.

Pipeline is created to execute the Notebook





Job runs in Azure Data bricks



Conclusion

The "Inventory Management System" project successfully leveraged powerful data processing tools such as PySpark and Azure services to analyze and derive meaningful insights from large datasets. By transforming raw data into structured fact and dimension tables, the system implemented a scalable and efficient data pipeline capable of supporting comprehensive analysis. The integration of **Delta Lake** and **Azure Databricks** enabled the execution of complex transformations and machine learning models at scale, delivering real-time analytics to enhance decision-making.

The incorporation of **Azure Data Factory** further streamlined the data processing and report generation workflows by automating tasks and minimizing manual intervention. This ensured the system was not only robust but also scalable to handle increasing data volumes effectively. The generated reports—spanning product sales metrics, customer trends, and transaction statistics—offered actionable insights into inventory and product performance. These insights can be leveraged to optimize sales tracking, improve inventory control, and inform strategic planning for future business growth.

Ultimately, this project demonstrated how the integration of big data tools and cloud services can transform raw data into actionable insights, enabling stakeholders to make data-driven decisions. By combining scalability, automation, and advanced analytics, the system empowers organizations to gain a deeper understanding of their inventory metrics and drive improved operational performance.

